

Ch9 Notes - 3/27/2023

Monday, March 27, 2023 10:54 AM

////////////////////////////////////

Lisa Jacklin

CS 321 Operating Systems

Lecture Notes 3/27/2023

////////////////////////////////////

Midterms are almost done being graded and hoping wednesday we will get them back

////////////////////////////////////

Chapter 9: Memory Management -> Virtual Memory

demand paging is our main focus today: demand paging is when an OS is demanding memory

page vs frame:

- page is in the virtual memory
- frame is the physical memory

page fault

- if there is a reference to a page, first reference to that page will trap to OS
- when the OS requests a page that does not exist
- when a process tries to find a page that is not in memory and must be swapped in from physical memory to be used without fault

->THIS IS IMPORTANT FOR QUIZ AND OR EXAM!<-

Steps in handling a page fault

1. reference page table,
2. trap OS
3. page is on the backing store

...

->the rest of the steps are on slide 16 of ch 9

performance of demanding paging

- the demand paging, the more time spend waiting for the operation to complete

stages:

- trap to the OS
- save the user registers and process state
- determine that the interrupt was a page fault
- check that the page reference was legal and determine the location of the page on the memory
- issue a read from the disk to a free frame
- while waiting, allocate the CPU to some other
- ... more on slide 19...

three major activities

- service the interrupt : far less time than reading the page
- read the page : takes a larger amount of time to read that data
- restart the process: only a small amount of time

page fault rate $0 \leq p \leq 1$

- if $p = 0$ no page faults
- if $p = 1$ every reference is at fault

Effective Access Time (EAT)

$$\text{EAT} = (1-p) \times \text{memory access} + p (\text{page fault overhead} + \text{swap page out} + \text{swap page in})$$

Demand Paging Example

- Memory access time = 200 nanoseconds
- Average page-fault service time = 8 milliseconds
- $\text{EAT} = (1 - p) \times 200 + p (8 \text{ milliseconds})$
- $= (1 - p) \times 200 + p \times 8,000,000$
- $= 200 + p \times 7,999,800$
- If one access out of 1,000 causes a page fault, then
- $\text{EAT} = 8.2 \text{ milliseconds}$
- This is a slowdown by a factor of 40!
- If want performance degradation < 10 percent:
- $220 > 200 + 7,999,800 \times p$
- $20 > 7,999,800 \times p$
- $p < .0000025$
- < one page fault in every 400,000 memory accesses

-> I would go through and make sure you understand the EAT equation!

- you can check your understanding with example on slide 21
- note that the average page-fault service time can include any of the three major activities listed above.

-> END OF MAJOR IMPORTANCE ON TESTS/QUIZES <-

Demand Paging Optimization (design considerations)

-> no questions on test over this

-> note that this uses SATA technology for some amount of time and now, faster....?

- Swap space I/O faster than file system I/O even if on the same device
 - o Swap allocates in larger chunks, less management needed than file system
- Copy entire process image to swap space at process load time
 - o Then page in and out of the swap space which is used primarily in older Unix systems
- Demand page in from program binary on disk but discard rather than paging out when freeing frame
 - o Still need to write to swap space
 - o Page not associated with a file (like stack and heap) - anonymous memory
 - o Pages modified in memory but not yet written back to the file system
- Mobile systems typically don't support
- swapping

Copy on Write

- Copy-on-Write: allows both parent and child processes to initially share the same pages in memory
 - o If either process modifies a shared page, only then is the page copied.
- COW allows more efficient process creation as only modified pages are copied.
- In general, free pages are allocated from a pool of zero fill on demand pages
 - o Pool should always have free frames for fast demand page execution
 - Don't want to have to free a frame as well as other processes on page fault
- Vfork() variation on fork() system call has parent suspend and child using copy-on-write address space of parent
 - o Designed to have child call exec()

- Very efficient

What happens if there is no free frame?

- Used up by process pages
- Also in demand from the kernel, I/O buffers, etc
- Page replacement: find some page in memory, but not really in use, page it out
 - Algorithm: terminate, swap out, replace?
 - Performance: want an algorithm which will result in minimum number of page faults

Page replacement

- Prevent over-allocation of memory by modifying page-fault service routine to include page replacement
- Use modify (dirty) bit to reduce overhead of page transfers only modified pages are written to disk
- Page replacement completes separation between logical memory and physical memory; large virtual memory can be provided on smaller memory

->view slide 29

Basic page replacement

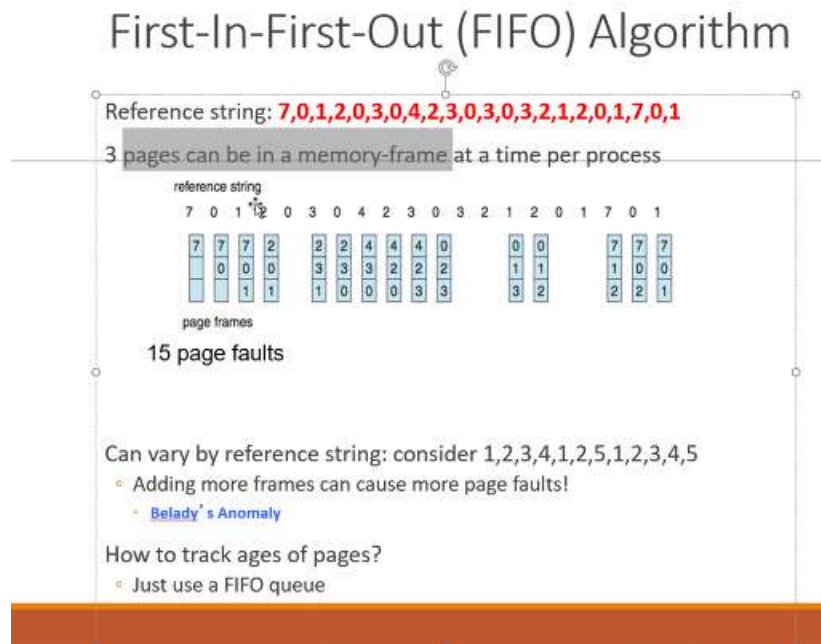
- Find the location of the desired page on the disk
- Find a free frame:
 - If there is one, use it
 - Algorithm selects the victim frame if there is no free frame
- Bring the desired page into the newly free frame; update the page and frame table
- Continue the process by restarting the instruction that caused the trap

...

Page and frame replacement algorithms

-> slide 31

- Primarily using reference string algorithm...



- The larger the size of the frame, the more page faults, the larger the number of frames, the less amount of page faults

-> 32, 33, 34 are all visual and count help with understanding if you take a look at them and take a

moment to understand them.. Have few notes but can explain if necessary ;)

Ended lecture on slide 35.